

# LearnerSpace: Handwritten Digit Recognition Using Convolutional Neural Networks (CNN)

Submitted by: Arjun Singh (23B1272)

July 3, 2025

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Dataset Overview</b>	<b>2</b>
<b>3</b>	<b>Model Architecture</b>	<b>2</b>
<b>4</b>	<b>Training and Validation</b>	<b>2</b>
4.1	Setup . . . . .	2
4.2	Training Results . . . . .	3
<b>5</b>	<b>Test Set Evaluation</b>	<b>3</b>
<b>6</b>	<b>Inference on Custom Image</b>	<b>3</b>
<b>7</b>	<b>Confusion Matrix</b>	<b>4</b>
<b>8</b>	<b>Performance Analysis</b>	<b>4</b>
8.1	Observations . . . . .	4
8.2	Expected vs Actual . . . . .	5
<b>9</b>	<b>Challenges and Improvements</b>	<b>5</b>
9.1	Challenges . . . . .	5
9.2	Proposed Improvements . . . . .	5
<b>10</b>	<b>Conclusion</b>	<b>5</b>

# 1 Introduction

Handwritten digit recognition is a classical machine learning problem and a common benchmark for evaluating computer vision techniques. This project focuses on implementing a Convolutional Neural Network (CNN) using PyTorch to classify digits from the MNIST dataset. The project workflow includes model creation, training, evaluation, and inference on external images.

## 2 Dataset Overview

The MNIST dataset contains 70,000 grayscale images of handwritten digits (60,000 for training and 10,000 for testing), each sized  $28 \times 28$  pixels. Each image is labeled with a digit from 0 to 9. Data is preprocessed using 'ToTensor()' transformation which normalizes the pixel values to the  $[0, 1]$  range.

## 3 Model Architecture

The CNN model used consists of the following layers:

- **Conv2D Layer 1:** 1 input channel, 64 output channels, kernel size 3
- **MaxPooling Layer 1:** kernel size 2
- **Conv2D Layer 2:** 64 input and output channels, kernel size 3
- **MaxPooling Layer 2:** kernel size 2
- **Conv2D Layer 3:** 64 input and output channels, kernel size 3
- **MaxPooling Layer 3:** kernel size 2
- **Flatten**
- **Fully Connected Layers:**  $[64 * 1 * 1 \rightarrow 64 \rightarrow 32 \rightarrow 10]$

The final layer outputs a vector of size 10 corresponding to the class scores for digits 0–9.

## 4 Training and Validation

### 4.1 Setup

- Optimizer: Adam
- Loss Function: CrossEntropyLoss
- Epochs: 5
- Batch Size: 64
- Train/Validation Split: 70/30

## 4.2 Training Results

Epoch	Loss	Train Accuracy	Validation Accuracy
1	0.718	76.04%	92.80%
2	0.255	94.00%	96.00%
3	0.191	96.30%	97.30%
4	0.159	97.14%	97.60%
5	0.130	97.90%	98.10%

Table 1: Training and Validation Accuracy per Epoch

## 5 Test Set Evaluation

The trained model was evaluated on the 10,000-sample MNIST test set.

- **Test Accuracy:** 98.35%

## 6 Inference on Custom Image

A custom handwritten digit image ('4.png') was loaded using OpenCV, converted to grayscale, resized to 28x28, normalized, and passed to the trained model.



Figure 1: Custom Digit Image (Displayed in RGB)

**Prediction:** 4

## 7 Confusion Matrix




Figure 2: Confusion Matrix on MNIST Test Set

The model performed well across all digit classes, with most misclassifications occurring between visually similar digits such as 4 and 9, and 5 and 3.

## 8 Performance Analysis

### 8.1 Observations

- The model achieved high accuracy on both validation and test sets, suggesting it generalizes well.
- Confusion matrix reveals most confusion between digits that are structurally similar.
- Performance is consistent and robust with minimal overfitting.

## 8.2 Expected vs Actual

The expected accuracy was around 97%, and the model surpassed it with a final accuracy of 98.35%. This aligns well with performance seen in standard CNN-based MNIST classifiers.

# 9 Challenges and Improvements

## 9.1 Challenges

- Ensuring the input dimensions match after each convolution and pooling layer.
- Handling external images which often differ in thickness, position, or background noise.

## 9.2 Proposed Improvements

1. **Data Augmentation:** Introduce transformations like rotation, zoom, and shift to make the model robust to variations in handwriting.
2. **Batch Normalization and Dropout:** Adding batch normalization layers can speed up training and stabilize the learning process, while dropout can prevent overfitting.

# 10 Conclusion

This project demonstrates the application of Convolutional Neural Networks in recognizing handwritten digits. The implemented model achieves over 98% accuracy on the MNIST dataset and performs well on external images. With further improvements such as data augmentation and regularization techniques, the model can become even more robust and production-ready.

## References

- PyTorch Documentation
- MNIST Dataset
- Scikit-learn Confusion Matrix